

Fuelling an Oil & Gas Giant

A diesel, haulage and SARS-refund data story — 23.7 million rows from Kalahari Petroleum's upstream fleet-operations ERP, modelled as a SQL Server star schema, shipped through Azure Data Factory, and served to Qlik Sense.

[View source on GitHub →](#)

Case-study note: Kalahari Petroleum is a fictional company invented for this portfolio piece. The underlying operational data is real (anonymised) mining/haulage fleet-fuel ERP data, presented here under a fictional oil & gas identity to demonstrate the data model and analytics pipeline without naming the source organisation.

290.6M L

fuel issued 2009–2022

325 504

AFS fuel transactions

778 254

equipment trips

R 385.2m

modelled diesel refunds

23.7M

rows through the pipeline


1 • More than fuel records

The source system is an operational ERP for Kalahari Petroleum's upstream fleet: an Automated Fuel System (AFS) issuing diesel to haul trucks and drill rigs, tank deliveries, odometer and hour-meter readings, haulage trips with material types, GPS trip traces over a 21.9-million-point coordinate grid, geofences, SAP integration staging and cost-centre accounting. Fuel is the connective tissue — every litre issued ties equipment, location, activity and money together. Kalahari Petroleum's fleet-management system classifies vehicles the way a mine would — drill rigs, haul trucks, waste removal — because well-pad

construction and haulage logistics on an oil & gas site mirror mining operations almost exactly.


This project models that data as a Kimball star schema (8 dimensions, 8 facts) with one high-stakes business calculation at its centre: the **SARS diesel refund** under Rebate Item 670.04 of the Customs & Excise Act, where classifying litres as eligible or non-eligible is worth millions of rand per month.

2 · The fuel story





Fuel issue volumes step up sharply from 2019 as the AFS rollout reaches full coverage, settling around 70–80 million litres a year across the operation.

Monthly fuel issued (litres, 2019 onward)



Fuel issued by depot / location (top 12)







Demand is not smooth: maintenance shutdowns, production cycles and season shape the curve — one reason a refund forecast needs more than a flat average.


3 · Fleet & haulage

Top 15 fuel-consuming equipment (red = SARS-ineligible class)




Fuel issued by vehicle type (top 12)






Trip records carry the material being moved — extracted product, waste, consumables — which is exactly the evidence SARS eligibility classification leans on: transport of extracted material on the production site is claimable, general road use is not.



4 · Haulage & site activity map

Real GPS trip data plotted against a label-free basemap — three toggleable layers built for three different audiences: **dispatch** (where the traffic actually is), **production** (what's moving out of each site), and **compliance** (which activity supports a SARS eligible-activity claim). The top 40 haul routes overlay shows road-maintenance priority by traffic volume.




45 haulage sites · 694,217 trips · real GPS coordinates (site/company identity kept fictional — see case-study note — so the basemap below is intentionally label-free: no place names, only terrain/road geometry for genuine ops/dispatch use). Hover a site for a quick summary, click for full detail. **Data-quality note:** the source system's Lat/Long columns are transposed (the "Lat" column holds longitude values and vice versa) — corrected here. Map requires an internet connection to load street tiles.

5 · The SARS diesel refund

For on-land primary producers (mining included), the refund formula from the policy evidence is:

```
eligible_litres    = total_litres - non_eligible_litres
qualifying_litres = eligible_litres × 80%
refund_rand       = qualifying_litres × refund_rate (c/L) ÷ 100
```




Last 12 claim months

Claim month	Total L	Non-eligible L	Eligible L	Qualifying L	Refund
2021-09	5 982 077	2 471 924	3 510 153	2 808 122	R 9 800 347.18
2021-10	6 455 776	2 405 818	4 049 958	3 239 966	R 11 307 482.74
2021-11	5 604 740	2 285 870	3 318 870	2 655 096	R 9 266 285.04
2021-12	5 922 270	2 197 315	3 724 955	2 979 964	R 10 400 075.48
2022-01	5 972 874	2 297 767	3 675 107	2 940 086	R 10 260 899.86
2022-02	5 612 826	2 345 327	3 267 500	2 614 000	R 9 122 859.44
2022-03	6 272 003	3 143 907	3 128 096	2 502 477	R 8 733 643.19
2022-04	6 193 891	2 914 930	3 278 961	2 623 169	R 9 154 858.83
2022-05	6 778 019	3 279 961	3 498 059	2 798 447	R 9 766 579.61
2022-06	6 344 588	3 212 236	3 132 352	2 505 881	R 8 745 525.95

2022-07	6 473 722	3 384 942	3 088 781	2 471 025	R 8 623 875.99
2022-08	13 224 056	10 952 680	2 271 376	1 817 101	R 6 341 680.95

Caution: rates are the 2020 SARS policy examples found in the project's evidence pack (349 c/L on-land after 1 Apr 2020). They are implementation evidence, not tax advice — update `dw.DimRefundRate` with current rates before relying on any figure.

6 · Data quality — where the money leaks




Finding	Rows
Usage row unmatched to equipment	203 309
Fuel issue exceeds 1.5x tank size	14 746
AFS duplicate FuelEventId	1 728
Delivery zero/negative volume	288

The stand-outs: over 200 000 usage-logbook rows carry registration numbers that don't match the equipment master (an audit-trail gap for refund claims), and ~14 700 single fuel issues exceed 1.5× the receiving vehicle's tank size — classic candidates for meter faults or leakage/theft review. One equipment-master row even had a tab character embedded in its


registration number, found because it was the only row in 23.7 million that broke a rectangular file export.

7 · Machine learning: fuel-anomaly detection

An **Isolation Forest** (scikit-learn, 300 trees, 2% contamination) scores every fuel transaction on five features: litres issued, tank-fill ratio, that vehicle's own fill-ratio z-score (so a naturally large tanker isn't penalised for being large), hour of day, and day of week. This catches what a fixed 1.5×-tank-size rule misses — multivariate outliers like a small bakkie fuelling at an odd hour for a ratio that's unremarkable on its own but anomalous in combination.



6 511 of 325 504 transactions flagged (2.0%). The highest-scoring cases are extreme: single transactions of hundreds of thousands to millions of litres against 100–600 L tanks — almost certainly decimal-placement data-entry errors rather than physical theft, but exactly the kind of finding a refund audit needs caught before litres get claimed.



Top 10 highest-risk transactions

Date/time	Fleet ID	Make	Litres	Tank (L)	Fill ratio	Anomaly score
11/08/2020 00:07:04	DB338	Diesel Bowser	30 202.0	500	60.4x	0.801
05/30/2019 01:01:03	MOG1475	TOYOTA	15 901.5	130	122.3x	0.800
05/29/2019 17:36:20	MOG1410	TOYOTA	159 015.1	150	1 060.1x	0.800
08/11/2022 11:52:33	CON685	CATERPILLAR	8 000 000.0	625	12 800.0x	0.799
02/17/2020 16:38:34	CON093	CATERPILLAR	21 436.0	532	40.3x	0.799
05/30/2019 06:24:46	MOG1396	TOYOTA	15 901.5	150	106.0x	0.799
09/26/2020 05:57:43	DB338	Diesel Bowser	28 659.5	500	57.3x	0.797

05/29/2019 22:08:25	MOG1420	TOYOTA	159 015.1	150	1 060.1x	0.797
05/30/2019 01:13:18	CNL349L	TOYOTA	15 901.5	80	198.8x	0.797
07/22/2019 02:02:38	BULK270	CATERPILLAR	32 096.6	750	42.8x	0.796

Full 200-row review queue: data/analysis/ml_review_queue.csv and the "ML Anomaly Review Queue" sheet in the Excel workbook. Model: sql/10_ml_anomaly_detection.py.

9 · How it's built

```
SQL Server (source ERP, QA copy)           — local, always works offline
├ dw star schema · 8 dims, 8 facts, 5 views (01_create_load_dw_full.sql)
├ TSV export (bcp, UTF-8)                  (02_export_dw_to_tsv.ps1)
│   ├── curated_local/ Parquet              (03_build_local_parquet.py)
│   └─ azcopy → ADLS Gen2 raw/              — Azure mirror
│       ├── Data Factory pl_raw_to_curated (ForEach Copy, TSV → snappy Parquet)
│       │   └─ curated/dw/<Table>/*.parquet
│       └─ Qlik Sense Cloud (SAS web-files or Azure Storage connector)
```

The warehouse reconciles exactly: 290 557 288.29 litres issued in both source and fact table, to the hundredth of a litre. Everything regenerates from five scripts, and the Azure resource group tears down with one command when it has served its purpose.

Full source, the SQL build script, both Qlik load scripts (cloud + local), and the Azure Data Factory definitions are in the repository:

https://github.com/anthonyapollis/KalahariPetroleum_DW_Azure_Qlik. See

[qlik/QLIK_APP_GUIDE.md](#) for the six-sheet Qlik Sense app design, and

[azure/AZURE_ARCHITECTURE.md](#) for the deployed-resource inventory and kill switch.

Anthony Apollis · 2026 · Built with SQL Server, Azure Data Factory, ADLS Gen2, Python & Qlik Sense.
QA data, anonymised context, presented under a fictional company; SARS figures are modelled examples, not tax advice.

[github.com/anthonyapollis/KalahariPetroleum DW Azure Qlik](https://github.com/anthonyapollis/KalahariPetroleum_DW_Azure_Qlik)